

Jithin George

www.linkedin.com/in/jithu010 | +919633473792 | jithin.george.w@gmail.com

SUMMARY

Backend Developer with 2+ years of experience building AI-first backend systems using Python, FastAPI, and modern vector databases. Proven expertise in deploying GPT-powered RAG pipelines, scalable API services, and generative AI tools using LangChain, LlamaIndex, and cloud-native infrastructure. MSc (AI) graduate with solid ML/DL foundations and real-world enterprise deployment experience.

SKILLS & INTERESTS

Technical Skills:	Python, JavaScript (Node.js), FastAPI, LangChain, LlamaIndex, REST APIs, Microservices, SQL, MongoDB, Weaviate, Qdrant, Elasticsearch, Neo4j, Docker, Git, AWS, Azure, MLOps, CI/CD, Vector Databases, Redis, Semantic Search, Scalability Optimization.
AI/ML Expertise:	GPT-3.5, GPT-4, Prompt Engineering, ReAct, Chain-of-Thought, Retrieval-Augmented Generation (RAG), Embedding-Based Search, Machine Learning, Deep Learning, scikit-learn, TensorFlow, PyTorch, Model Monitoring, Model Evaluation, AIOps Pipelines, Ethical AI Practices.

WORK EXPERIENCE

Gapblue Software Labs <i>Associate Consultant,</i>	Ernakulam, Kerala <i>Mar 2024 - Present</i>
- Translated complex business problems into scalable LLM solutions using LangChain, LlamaIndex, and Azure OpenAI, improving internal semantic search speed by 40% and retrieval accuracy by 35%. - Presented model outputs and performance evaluations to cross-functional teams, enabling data-driven decisions that improved product design and usability. - Led architecture of <i>Knowledge Explorer</i>, a multilingual, glossary-aware GPT-4 assistant integrated with Azure OpenAI and Translator, boosting internal knowledge accessibility by 60%. - Designed and deployed vector database integrations (Qdrant, Weaviate, Pinecone), reducing lookup latency 3× and maintaining sub-300ms response times for real-time AI use cases. - Delivered a successful client-facing RAG Q&A solution that won a \$500K contract, contributing to a 22% increase in company revenue within 6 months. - Automated document-heavy workflows across 4 departments via RAG pipelines, cutting manual effort by 50%, saving an estimated \$200K annually. - Optimized GPT-3.5/4o inference pipelines to reduce latency by 35% and API cost by 20%, leading to \$250K in operational savings. - Worked closely with product and engineering teams in Agile sprints to deliver 3+ production-grade GenAI tools with 30% higher client satisfaction than previous benchmarks.	
<i>Trainee Associate Applications Engineer,</i>	<i>Sep 2023 – Feb 2024</i>
- Supported backend and full-stack development across Node.js, FastAPI, and React, building foundational production-ready features. - Enabled smoother GPT-4 integration into RAG-based tools, cutting inference latency by 25% and accelerating internal search performance. - Collaborated on a research and patent assistant bot, transforming complex technical documents into searchable AI queries for faster innovation access. - Participated in bug resolution sprints and cross-functional stand-ups, reducing debugging cycles by 40%.	

EDUCATION

MSc Computer Science (AI), CUSAT, Kochi <i>Project: Forecasted AQI & heatwaves using Autoencoder-ARIMA + BiLSTM/DNN hybrid model.</i>	2021–2023
BCA, Rajagiri College, Kochi <i>Project: E-commerce web app with user auth & order management (HTML, JS, PHP, SQL).</i>	2018–2021

ADDITIONAL INFORMATION

Certifications: Oracle AI Vector Search Professional, OCI AI Foundations Associate (2023).