

Kevin Cherian George

+91 8078776093 | kevin2008.work@gmail.com | linkedin.com/in/kevin-cherian-george | github.com/kevin-291

EDUCATION

Karunya Institute of Technology and Sciences

Bachelor of Technology in Electronics and Computer Engineering; CGPA: 8.35/10

Coimbatore, India

July 2021 – July 2025

EXPERIENCE

AI/ML Engineer

Akiyam Solutions Private Limited

Aug. 2025 – Present

Vadodara, India

- Built multi-agent AI systems with LangGraph and Groq APIs orchestrating 6 specialized agents for simulation generation, plus ensemble vision pipelines (YOLOv8 + DETR) with weighted IOU fusion improving document layout detection accuracy through confidence-based model consensus.
- Optimized ML infrastructure implementing performance telemetry (token/latency/cost tracking), dynamic resource management (GPU-to-CPU fallback, adaptive batching) for large-scale OCR pipelines, reducing CUDA OOM errors and enabling data-driven optimization across distributed LLM operations.
- Delivered production AI frameworks automating workflows from NLP parsing to validated code/document outputs, integrating Hugging Face models, custom runners, parallel processing, and automated evaluation metrics for enterprise-scale simulation generation and document digitization systems.

Computer Vision Intern

Fraunhofer IDMT

Sep. 2024 – Nov. 2024

Ilmenau, Germany

- Designed and developed a video scraping pipeline to automatically extract and organize broadcast media from various German news outlets—reduced dataset collection time by more than 80 %
- Enhanced face tracking and recognition performance by integrating a modified face detection deep learning model, improving face detection accuracy on high-res TV media by an estimated 15 % (based on internal benchmark comparisons)
- Built a real-time inference pipeline using Python, OpenCV, ONNX Runtime, NumPy, and Pandas for extracting detection regions and embeddings from video frames—enabled high-throughput media analysis

PROJECTS

PoxViT | Python, PyTorch, NumPy, Scikit-Learn, Docker

- Designed and implemented PoxViT, a novel Vision Transformer architecture for skin disease classification, incorporating convolutional patch embedding for enhanced spatial feature extraction
- Developed end-to-end PyTorch pipeline with customizable architecture parameters
- Generated and analyzed attention maps for model interpretability
- Achieved an accuracy of 98 %

FinSight AI | Python, Pydantic AI, Asyncio, Ollama, OpenAI API

- Developed multi-agent AI pipeline extracting entities, analyzing sentiment, and predicting market impact from financial news
- Architected three specialized agents using Pydantic AI with local LLM deployment via Ollama (gemma3:1b model)
- Implemented async batch processing with asyncio for concurrent analysis of multiple articles, improving efficiency by 3x
- Built custom evaluation framework with Pydantic schemas validating outputs against ground truth with 85 %+ accuracy
- Designed JSON extraction system converting unstructured news into structured data (tickers, investments, sector info)

TECHNICAL SKILLS

Languages: Python, Java, C/C++, MySQL, JavaScript, HTML/CSS

Developer Tools: Git, Docker, GCP, AWS, Azure, VS Code, Visual Studio, PyCharm, Roboflow, MLflow, Open AI API, Groq API, Langchain, Langgraph

Libraries: Pandas, NumPy, Matplotlib, Scikit-Learn, PyTorch, Tensorflow, Keras, OpenCV, ONNX Runtime